

Double-click (or enter) to edit

## ▼ Project #1: Bigmart Sale Prediction

### Initializing Packages and Importing Data

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import warnings # Ignores any warning
warnings.filterwarnings("ignore")

train = pd.read_csv("data/Train.csv")
test = pd.read_csv("data/Test.csv")
```

### Taking a peak at our data

```
train.head()
```

	Item_Identifier	Item_Weight	Item_Fat_Content	Item_Visibility	Item_Type	Item_
0	FDA15	9.30	Low Fat	0.016047	Dairy	249.8
1	DRC01	5.92	Regular	0.019278	Soft Drinks	48.2
2	FDN15	17.50	Low Fat	0.016760	Meat	141.6
3	FDX07	19.20	Regular	0.000000	Fruits and Vegetables	182.0
4	NCD19	8.93	Low Fat	0.000000	Household	53.8

```
train.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 8523 entries, 0 to 8522
Data columns (total 12 columns):
```

2/4/22, 2:19 PM

Bigmart Sales Data Set.ipynb - Colaboratory

#	Column	Non-Null	Count	Dtype
0	Item_Identifier	8523	non-null	object
1	Item_Weight	7060	non-null	float64
2	Item_Fat_Content	8523	non-null	object
3	Item_Visibility	8523	non-null	float64
4	Item_Type	8523	non-null	object
5	Item_MRP	8523	non-null	float64
6	Outlet_Identifier	8523	non-null	object
7	Outlet_Establishment_Year	8523	non-null	int64
8	Outlet_Size	6113	non-null	object
9	Outlet_Location_Type	8523	non-null	object
10	Outlet_Type	8523	non-null	object
11	Item_Outlet_Sales	8523	non-null	float64

dtypes: float64(4), int64(1), object(7)  
memory usage: 799.2+ KB

train.describe()

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishment_Year	Item_Outlet_Sales
count	7060.000000	8523.000000	8523.000000	8523.000000	8
mean	12.857645	0.066132	140.992782	1997.831867	2
std	4.643456	0.051598	62.275067	8.371760	1
min	4.555000	0.000000	31.290000	1985.000000	
25%	8.773750	0.026989	93.826500	1987.000000	
50%	12.600000	0.053931	143.012800	1999.000000	1
75%	16.850000	0.094585	185.643700	2004.000000	3
max	21.350000	0.328391	266.888400	2009.000000	13

#Check for duplicates  
idsUnique = len(set(train.Item\_Identifier))  
idsTotal = train.shape[0]  
idsDupli = idsTotal - idsUnique  
print("There are " + str(idsDupli) + " duplicate IDs for " + str(idsTotal) + " total entries")  
  
There are 6964 duplicate IDs for 8523 total entries

1. Exploratory Data Analysis (EDA)


1.1. Univariate Distribution

1.1.1. Distribution of the target variable : Item\_Outlet\_Sales

plt.style.use('fivethirtyeight')  
plt.figure(figsize=(12,7))  
sns.distplot(train.Item\_Outlet\_Sales, bins = 25)  
plt.ticklabel\_format(style='plain', axis='x', scilimits=(0,1))

```
plt.xlabel("Item_Outlet_Sales")
plt.ylabel("Number of Sales")
plt.title("Item_Outlet_Sales Distribution")
```

```
Text(0.5, 1.0, 'Item_Outlet_Sales Distribution')
```



```
print ("Skew is:", train.Item_Outlet_Sales.skew())
print("Kurtosis: %f" % train.Item_Outlet_Sales.kurt())
```

```
Skew is: 1.1775306028542796
Kurtosis: 1.615877
```

### 1.1.2. Numerical Variables

```
numeric_features = train.select_dtypes(include=[np.number])
numeric_features.dtypes
```

```
Item_Weight          float64
Item_Visibility       float64
Item_MRP             float64
Outlet_Establishment_Year  int64
Item_Outlet_Sales     float64
dtype: object
```

```
numeric_features.corr()
```

	Item_Weight	Item_Visibility	Item_MRP	Outlet_Establishme
Item_Weight	1.000000	-0.014048	0.027141	-(
Item_Visibility	-0.014048	1.000000	-0.001315	-(
Item_MRP	0.027141	-0.001315	1.000000	(
Outlet_Establishment_Year	-0.011588	-0.074834	0.005020	1
Item_Outlet_Sales	0.014123	-0.128625	0.567574	-(

```
corr = numeric_features.corr()
```

```
print (corr['Item_Outlet_Sales'].sort_values(ascending=False))
```

```
Item_Outlet_Sales      1.000000
Item_MRP               0.567574
Item_Weight            0.014123
Outlet_Establishment_Year -0.049135
Item_Visibility        -0.128625
Name: Item_Outlet_Sales, dtype: float64
```

```
#correlation matrix
```

```
f, ax = plt.subplots(figsize=(12, 9))
```

```
sns.heatmap(corr, vmax=.8, square=True);
```




### 1.1.3. Categorical Variables

#### 1.1.3.1. Distribution of the Item\_Fat\_Content



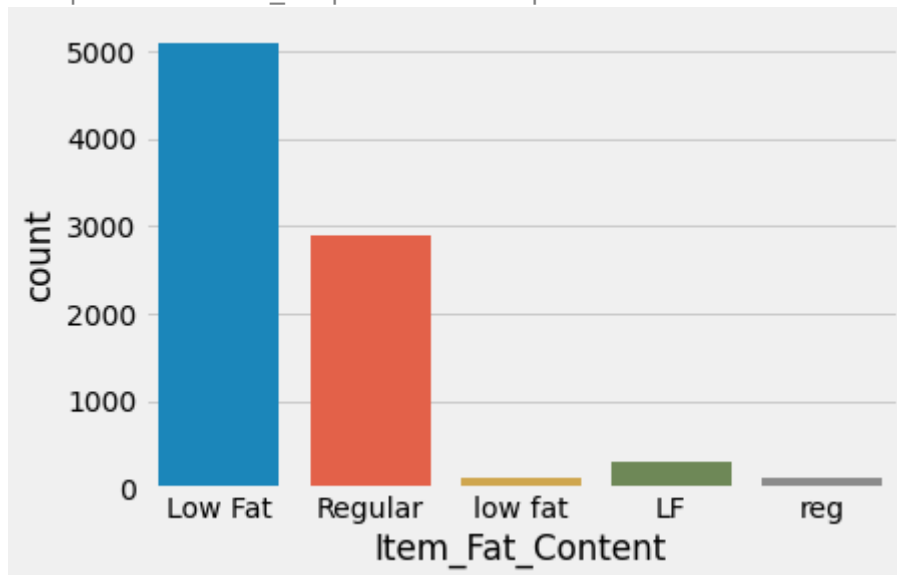
```
train.Item_Fat_Content.value_counts()
```

```
Low Fat    5089
Regular    2889
LF          316
reg         117
low fat     112
Name: Item_Fat_Content, dtype: int64
```



```
sns.countplot(train.Item_Fat_Content)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1c25668a10>
```



#### 1.1.3.2. Distribution of the Item\_Type

```
train.Item_Type.value_counts()
```

```
Fruits and Vegetables    1232
Snack Foods               1200
Household                 910
```

```

Frozen Foods      856
Dairy             682
Canned            649
Baking Goods     648
Health and Hygiene 520
Soft Drinks      445
Meat             425
Breads           251
Hard Drinks      214
Others           169
Starchy Foods   148
Breakfast        110
Seafood          64
Name: Item_Type, dtype: int64

```

```

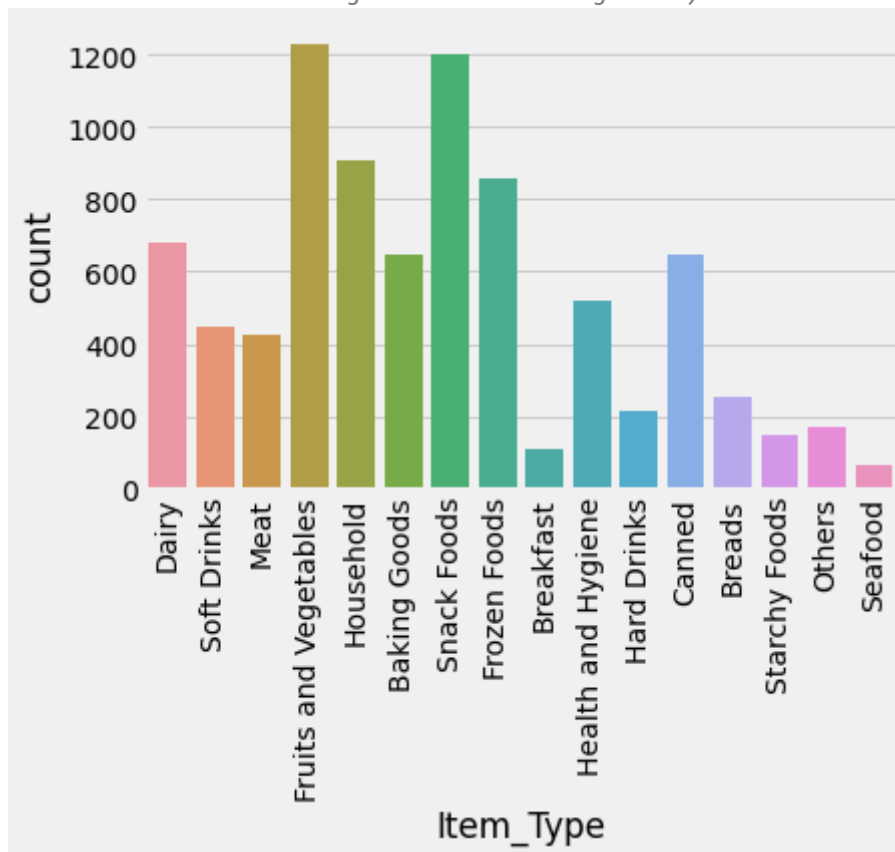
sns.countplot(train.Item_Type)
plt.xticks(rotation=90)

```

```

(array([ 0,  1,  2,  3,  4,  5,  6,  7,  8,  9, 10, 11, 12, 13, 14, 15]),
<a list of 16 Text major ticklabel objects>)

```



### 1.1.3.3. Distribution of the Outlet\_Size

```
train.Outlet_Size.value_counts()
```


```

Medium    2793
Small     2388
High       932
Name: Outlet_Size, dtype: int64

```

```
sns.countplot(train.Outlet_Size)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1c255e4810>
```




#### 1.1.3.4. Distribution of the Outlet\_Location\_Type

```
train.Outlet_Location_Type.value_counts()
```

```
Tier 3    3350
Tier 2    2785
Tier 1    2388
Name: Outlet_Location_Type, dtype: int64
```

```
sns.countplot(train.Outlet_Location_Type)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1c254efd50>
```



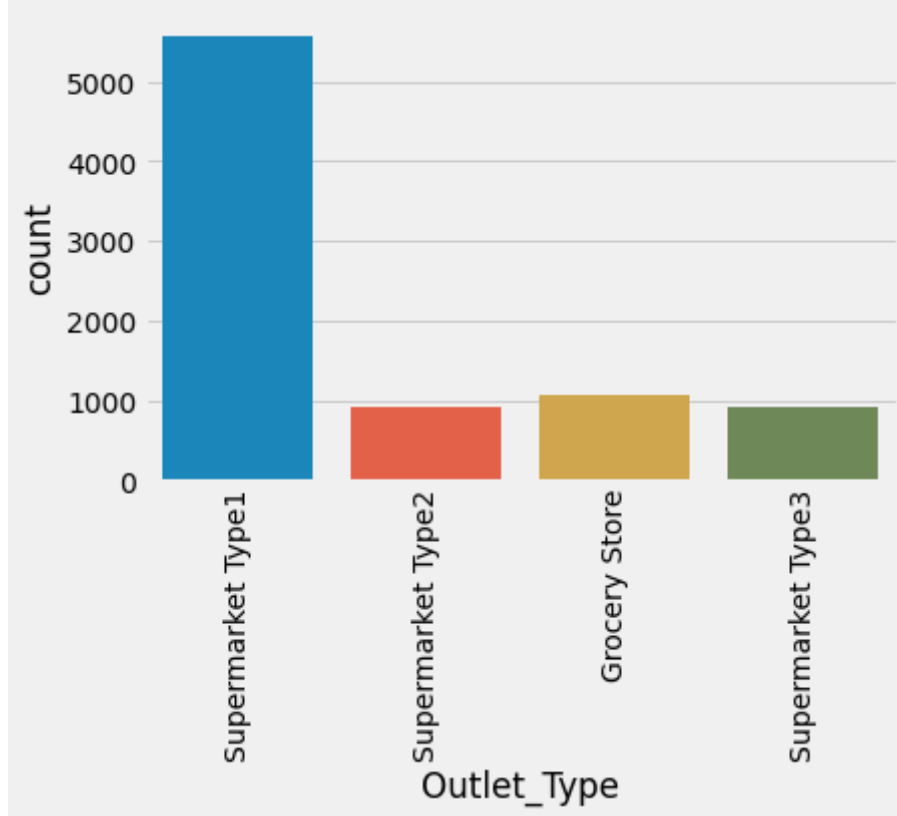
#### 1.1.3.5. Distribution of the Outlet\_Type

```
train.Outlet_Type.value_counts()
```

```
Supermarket Type1    5577  
Grocery Store        1083  
Supermarket Type3     935  
Supermarket Type2     928  
Name: Outlet_Type, dtype: int64
```

```
sns.countplot(train.Outlet_Type)  
plt.xticks(rotation=90)
```

```
(array([0, 1, 2, 3]), <a list of 4 Text major ticklabel objects>)
```



## 1.2. Bivariate Distribution


### 1.2.1. Numerical Variables

#### 1.2.1.1. Item\_Weight and Item\_Outlet\_Sales Analysis

```
plt.figure(figsize=(12,7))  
plt.xlabel("Item_Weight")  
plt.ylabel("Item_Outlet_Sales")  
plt.title("Item_Weight and Item_Outlet_Sales Analysis")  
plt.plot(train.Item_Weight, train["Item_Outlet_Sales"], '.', alpha = 0.3)
```




```
[<matplotlib.lines.Line2D at 0x7f1c253f1b90>]
```



### 1.2.1.2. Item\_Visibility and Item\_Outlet\_Sales Analysis

```
plt.figure(figsize=(12,7))
plt.xlabel("Item_Visibility")
plt.ylabel("Item_Outlet_Sales")
plt.title("Item_Visibility and Item_Outlet_Sales Analysis")
plt.plot(train.Item_Visibility, train["Item_Outlet_Sales"],'.', alpha = 0.3)
```


```
[<matplotlib.lines.Line2D at 0x7f1c23b68a10>]
```



### 1.2.1.3. Item\_MRP and Item\_Outlet\_Sales Analysis

```
plt.figure(figsize=(12,7))
plt.xlabel("Item_MRP")
plt.ylabel("Item_Outlet_Sales")
plt.title("Item_MRP and Item_Outlet_Sales Analysis")
plt.plot(train.Item_MRP, train["Item_Outlet_Sales"],'.', alpha = 0.3)
```


[<matplotlib.lines.Line2D at 0x7f1c23ae0410>]



### 1.2.1.4. Outlet\_Establishment\_Year and Item\_Outlet\_Sales Analysis

```
Outlet_Establishment_Year_pivot = \
train.pivot_table(index='Outlet_Establishment_Year', values="Item_Outlet_Sales", aggfunc=n

Outlet_Establishment_Year_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Outlet_Establishment_Year")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Outlet_Establishment_Year on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```



## 1.2.2. Categorical Variables


```
numeric_features = train.select_dtypes(include=[np.object])
numeric_features.dtypes
```

```
Item_Identifier      object
Item_Fat_Content     object
Item_Type            object
Outlet_Identifier    object
Outlet_Size          object
Outlet_Location_Type object
Outlet_Type          object
dtype: object
```

### 1.2.2.1. Impact of Item\_Identifier on Item\_Outlet\_Sales

```
plt.figure(figsize=(12,7))
plt.xlabel("Item_Identifier")
plt.ylabel("SQRT Item_Outlet_Sales")
plt.title("Item_Identifier and Item_Outlet_Sales Analysis")
plt.plot(train.Item_Identifier , train["Item_Outlet_Sales"],'.', alpha = 0.3)
```


[&lt;matplotlib.lines.Line2D at 0x7f1c2567c210&gt;]



### 1.2.2.2. Impact of Item\_Fat\_Content on Item\_Outlet\_Sales

```
Item_Fat_Content_pivot = \
train.pivot_table(index='Item_Fat_Content', values="Item_Outlet_Sales", aggfunc=np.median)

Item_Fat_Content_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Item_Fat_Content")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Item_Fat_Content on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```




1.2.2.3. Impact of Outlet\_Identifier on Item\_Outlet\_Sales



```
Outlet_Identifier_pivot = \
train.pivot_table(index='Outlet_Identifier', values="Item_Outlet_Sales", aggfunc=np.median)

Outlet_Identifier_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Outlet_Identifier ")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Outlet_Identifier on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```



```
train.pivot_table(values='Outlet_Type', columns='Outlet_Identifier',aggfunc=lambda x:x.mod
```

Outlet_Identifier	OUT010	OUT013	OUT017	OUT018	OUT019	OUT027
Outlet_Type	Grocery Store	Supermarket Type1	Supermarket Type1	Supermarket Type2	Grocery Store	Supermarket Type3


```
train.pivot_table(values='Outlet_Type', columns='Outlet_Size',aggfunc=lambda x:x.mode())
```

Outlet_Size	High	Medium	Small
Outlet_Type	Supermarket Type1	Supermarket Type3	Supermarket Type1

#### 1.2.2.4. Impact of Outlet\_Size on Item\_Outlet\_Sales

```
Outlet_Size_pivot = \
train.pivot_table(index='Outlet_Size', values="Item_Outlet_Sales", aggfunc=np.median)
```

```
Outlet_Size_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Outlet_Size ")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Outlet_Size on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```




#### 1.2.2.5. Impact of Outlet\_Location\_Type on Item\_Outlet\_Sales

```
Outlet_Location_Type_pivot = \
train.pivot_table(index='Outlet_Location_Type', values="Item_Outlet_Sales", aggfunc=np.med
```

```
Outlet_Location_Type_pivot.plot(kind='bar', color='blue',figsize=(12,7))
```

```
plt.xlabel("Outlet_Location_Type ")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Outlet_Location_Type on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```




```
train.pivot_table(values='Outlet_Location_Type', columns='Outlet_Type',aggfunc=lambda x:x.
```

Outlet_Type	Grocery Store	Supermarket Type1	Supermarket Type2	Supermarket Type3
Outlet_Location_Type	Tier 3	Tier 2	Tier 3	Tier 3

1.2.2.6. Impact of Outlet\_Type on Item\_Outlet\_Sales

```
Outlet_Type_pivot = \
train.pivot_table(index='Outlet_Type', values="Item_Outlet_Sales", aggfunc=np.median)

Outlet_Type_pivot.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Outlet_Type ")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Outlet_Type on Item_Outlet_Sales")
plt.xticks(rotation=0)
plt.show()
```

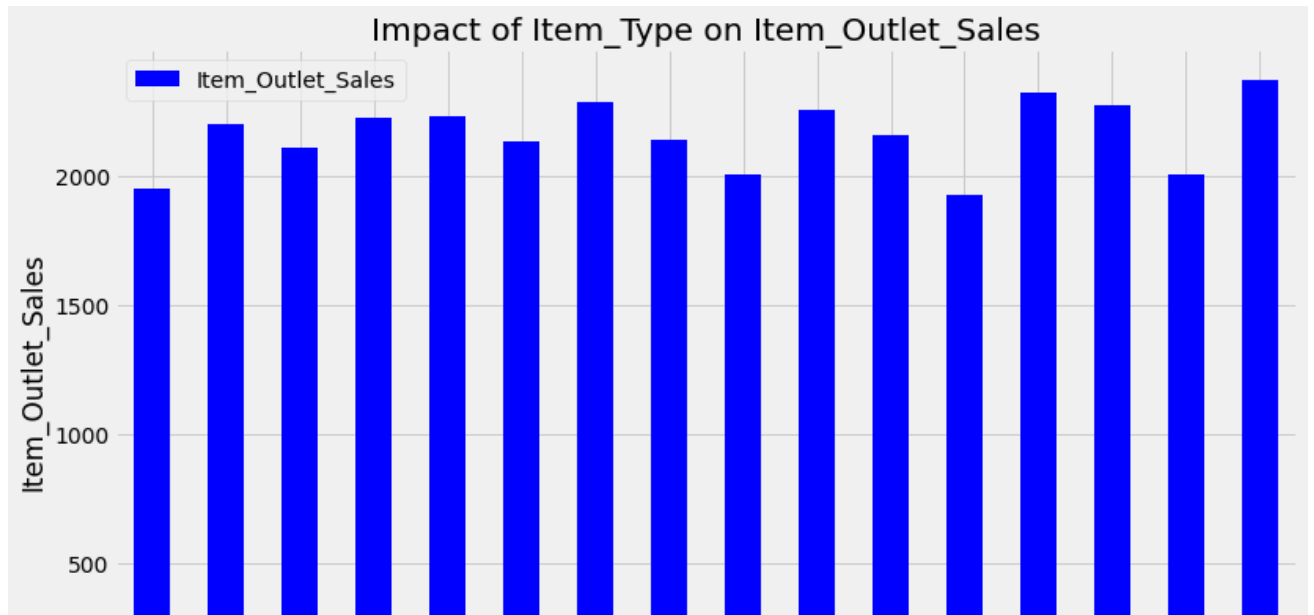


#### 1.2.2.7. Impact of Item\_Type on Item\_Outlet\_Sales

```
pivoTable = \
train.pivot_table(index='Item_Type', values="Item_Outlet_Sales", aggfunc=np.mean)

pivoTable.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Item_Type ")
plt.ylabel("Item_Outlet_Sales")
plt.title("Impact of Item_Type on Item_Outlet_Sales")
plt.xticks(rotation=90)
plt.show()
```






### 1.2.2.8. Impact of Item\_Type vs Item\_Visibility

```

od  ad  as  re  fir  od  le  nk  an  ok  ea  er  po  od  nk  od
pivoTable = \
train.pivot_table(index='Item_Type', values="Item_Visibility", aggfunc=np.mean)

pivoTable.plot(kind='bar', color='blue',figsize=(12,7))
plt.xlabel("Item_Type ")
plt.ylabel("Item_Visibility")
plt.title("Item_Type vs Item_Visibility")
plt.xticks(rotation=90)
plt.show()

```



## 2. Data Pre-Processing

### 2.1. Looking for missing values



```
# Join Train and Test Dataset
train['source']='train'
test['source']='test'

data = pd.concat([train,test], ignore_index = True)
data.to_csv("data/data.csv",index=False)
print(train.shape, test.shape, data.shape)

(8523, 13) (5681, 12) (14204, 13)
```



### 2.2. Imputing Missing Values

```
#aggfunc is mean by default! Ignores NA by default
item_avg_weight = data.pivot_table(values='Item_Weight', index='Item_Identifier')
print(item_avg_weight)
```

Item_Identifier	Item_Weight
DRA12	11.600
DRA24	19.350
DRA59	8.270
DRB01	7.390
DRB13	6.115
...	...
NCZ30	6.590
NCZ41	19.850
NCZ42	10.500
NCZ53	9.600
NCZ54	14.650

```
[1559 rows x 1 columns]
```

```
def impute_weight(cols):
    Weight = cols[0]
    Identifier = cols[1]

    if pd.isnull(Weight):
        return item_avg_weight['Item_Weight'][item_avg_weight.index == Identifier]
```

```

else:
    return Weight

```

```

print ('Original #missing: %d'%sum(data['Item_Weight'].isnull()))
data['Item_Weight'] = data[['Item_Weight','Item_Identifier']].apply(impute_weight,axis=1).
print ('Final #missing: %d'%sum(data['Item_Weight'].isnull()))

```

```

Original #missing: 2439
Final #missing: 0

```

## 2.3. Imputing Outlet\_size with the mode

```

#Import mode function:
from scipy.stats import mode

```

```

#Determing the mode for each
outlet_size_mode = data.pivot_table(values='Outlet_Size', columns='Outlet_Type',aggfunc=la
outlet_size_mode

```

Outlet_Type	Grocery Store	Supermarket Type1	Supermarket Type2	Supermarket Type3
Outlet_Size	Small	Small	Medium	Medium

```

def impute_size_mode(cols):
    Size = cols[0]
    Type = cols[1]
    if pd.isnull(Size):
        return outlet_size_mode.loc['Outlet_Size'][outlet_size_mode.columns == Type][0]
    else:
        return Size

```

```

print ('Original #missing: %d'%sum(data['Outlet_Size'].isnull()))
data['Outlet_Size'] = data[['Outlet_Size','Outlet_Type']].apply(impute_size_mode,axis=1)
print ('Final #missing: %d'%sum(data['Outlet_Size'].isnull()))

```

```

Original #missing: 4016
Final #missing: 0

```

## 3. Feature Engineering

### 3.1. Should we combine Outlet\_Type?

```

#Creates pivot table with Outlet_Type and the mean of Item_Outlet_Sales. Agg function is b
data.pivot_table(values='Item_Outlet_Sales', columns='Outlet_Type')

```

Outlet_Type	Grocery Store	Supermarket Tvne1	Supermarket Tvne2	Supermarket Tvne3
-------------	------------------	----------------------	----------------------	----------------------

### 3.2. Item\_Visibility minimum value 0

```
#Get all Item_Visibility mean values for respective Item_Identifier
visibility_item_avg = data.pivot_table(values='Item_Visibility',index='Item_Identifier')

def impute_visibility_mean(cols):
    visibility = cols[0]
    item = cols[1]
    if visibility == 0:
        return visibility_item_avg['Item_Visibility'][visibility_item_avg.index == item]
    else:
        return visibility

print ('Original #zeros: %d'%sum(data['Item_Visibility'] == 0))
data['Item_Visibility'] = data[['Item_Visibility','Item_Identifier']].apply(impute_visibil
print ('Final #zeros: %d'%sum(data['Item_Visibility'] == 0))

Original #zeros: 879
Final #zeros: 0
```

### 3.3. Determine the years of operation of a store

```
#Years:
data['Outlet_Years'] = 2013 - data['Outlet_Establishment_Year']
data['Outlet_Years'].describe()

count      14204.000000
mean         15.169319
std          8.371664
min           4.000000
25%          9.000000
50%         14.000000
75%         26.000000
max         28.000000
Name: Outlet_Years, dtype: float64
```

### 3.4. Create a broad category of Type of Item

```
#Get the first two characters of ID:
data['Item_Type_Combined'] = data['Item_Identifier'].apply(lambda x: x[0:2])
#Rename them to more intuitive categories:
data['Item_Type_Combined'] = data['Item_Type_Combined'].map({'FD':'Food',
                                                             'NC':'Non-Consumable',
                                                             'DR':'Drinks'})

data['Item_Type_Combined'].value_counts()

Food      10201
```

```

Non-Consumable      2686
Drinks              1317
Name: Item_Type_Combined, dtype: int64

```

### 3.5. Modify categories of Item\_Fat\_Content

```

#Change categories of low fat:
print('Original Categories:')
print(data['Item_Fat_Content'].value_counts())

print('\nModified Categories:')
data['Item_Fat_Content'] = data['Item_Fat_Content'].replace({'LF':'Low Fat',
                                                            'reg':'Regular',
                                                            'low fat':'Low Fat'})

print(data['Item_Fat_Content'].value_counts())

```

```

Original Categories:
Low Fat      8485
Regular      4824
LF           522
reg          195
low fat      178
Name: Item_Fat_Content, dtype: int64

```

```

Modified Categories:
Low Fat      9185
Regular      5019
Name: Item_Fat_Content, dtype: int64

```

```

#Mark non-consumables as separate category in low_fat:
data.loc[data['Item_Type_Combined']=="Non-Consumable", 'Item_Fat_Content'] = "Non-Edible"
data['Item_Fat_Content'].value_counts()

```

```

Low Fat      6499
Regular      5019
Non-Edible    2686
Name: Item_Fat_Content, dtype: int64

```

## 4. Feature Transformations

### 4.1. Creating variable Item\_Visibility\_MeanRatio

```

func = lambda x: x['Item_Visibility']/visibility_item_avg['Item_Visibility'][visibility_it
data['Item_Visibility_MeanRatio'] = data.apply(func,axis=1).astype(float)
data['Item_Visibility_MeanRatio'].describe()

```

```

count      14204.000000
mean        1.061884
std         0.235907

```

```

min            0.844563
25%            0.925131
50%            0.999070
75%            1.042007
max            3.010094
Name: Item_Visibility_MeanRatio, dtype: float64

```

## 4.2. Numerical and Categorical Variables – Dummy variables

```

#Import library:
from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()

#New variable for outlet
data['Outlet'] = le.fit_transform(data['Outlet_Identifier'])
var_mod = ['Item_Fat_Content', 'Outlet_Location_Type', 'Outlet_Size', 'Item_Type_Combined', '0']
le = LabelEncoder()
for i in var_mod:
    data[i] = le.fit_transform(data[i])

#Dummy Variables:
data = pd.get_dummies(data, columns=['Item_Fat_Content', 'Outlet_Location_Type', 'Outlet_Siz
    'Item_Type_Combined', 'Outlet'])

```

```
data.dtypes
```

```

Item_Identifier      object
Item_Weight          float64
Item_Visibility      float64
Item_Type            object
Item_MRP             float64
Outlet_Identifier     object
Outlet_Establishment_Year  int64
Item_Outlet_Sales    float64
source              object
Outlet_Years         int64
Item_Visibility_MeanRatio float64
Item_Fat_Content_0   uint8
Item_Fat_Content_1   uint8
Item_Fat_Content_2   uint8
Outlet_Location_Type_0 uint8
Outlet_Location_Type_1 uint8
Outlet_Location_Type_2 uint8
Outlet_Size_0        uint8
Outlet_Size_1        uint8
Outlet_Size_2        uint8
Outlet_Type_0        uint8
Outlet_Type_1        uint8
Outlet_Type_2        uint8
Outlet_Type_3        uint8
Item_Type_Combined_0 uint8
Item_Type_Combined_1 uint8
Item_Type_Combined_2 uint8
Outlet_0             uint8
Outlet_1             uint8
Outlet_2             uint8

```

```

Outlet_3      uint8
Outlet_4      uint8
Outlet_5      uint8
Outlet_6      uint8
Outlet_7      uint8
Outlet_8      uint8
Outlet_9      uint8
dtype: object

```

### 4.3. Exporting Data

```

#Drop the columns which have been converted to different types:
data.drop(['Item_Type', 'Outlet_Establishment_Year'],axis=1,inplace=True)

#Divide into test and train:
train = data.loc[data['source']=="train"]
test = data.loc[data['source']=="test"]

#Drop unnecessary columns:
test.drop(['Item_Outlet_Sales', 'source'],axis=1,inplace=True)
train.drop(['source'],axis=1,inplace=True)

#Export files as modified versions:
train.to_csv("data/train_modified.csv",index=False)
test.to_csv("data/test_modified.csv",index=False)

```

## 5. Model, predict and solve the problem

```

train_df = pd.read_csv('data/train_modified.csv')
test_df = pd.read_csv('data/test_modified.csv')

#Define target and ID columns:
target = 'Item_Outlet_Sales'
IDcol = ['Item_Identifier', 'Outlet_Identifier']
from sklearn import metrics
from sklearn.model_selection import cross_val_score

def modelfit(alg, dtrain, dtest, predictors, target, IDcol, filename):
    #Fit the algorithm on the data
    alg.fit(dtrain[predictors], dtrain[target])

    #Predict training set:
    dtrain_predictions = alg.predict(dtrain[predictors])

    #Perform cross-validation:
    cv_score = cross_val_score(alg, dtrain[predictors],(dtrain[target]) , cv=20, scoring='
    cv_score = np.sqrt(np.abs(cv_score))

    #Print model report:
    print("\nModel Report")

```

```
print("RMSE : %.4g" % np.sqrt(metrics.mean_squared_error((dtrain[target]).values, dtrain_predictions)))
print("CV Score : Mean - %.4g | Std - %.4g | Min - %.4g | Max - %.4g" % (np.mean(cv_scores), np.std(cv_scores),
                                np.min(cv_scores), np.max(cv_scores)))

#Print r2 score:
print("r2 score : %.4g" % metrics.r2_score((dtrain[target]).values, dtrain_predictions))

#Predict on testing data:
dtest[target] = alg.predict(dtest[predictors])

#Export submission file:
IDcol.append(target)
submission = pd.DataFrame({ x: dtest[x] for x in IDcol})
submission.to_csv(filename, index=False)
```

## Linear Regression Model

```
from sklearn.linear_model import LinearRegression
LR = LinearRegression(normalize=True)

predictors = train_df.columns.drop(['Item_Outlet_Sales', 'Item_Identifier', 'Outlet_Identifier'])
model = LinearRegression()
model.fit(LR, train_df, test_df, predictors, target, IDcol, 'LR.csv')

coef1 = pd.Series(LR.coef_, predictors).sort_values()
coef1.plot(kind='bar', title='Model Coefficients')
```



Model Report

RMSE : 1129

## Ridge Regression Model

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1c15df8c50>
```

```
from sklearn.linear_model import Ridge
RR = Ridge(alpha=0.05,normalize=True)
modelfit(RR, train_df, test_df, predictors, target, IDcol, 'RR.csv')
```

```
coef2 = pd.Series(RR.coef_, predictors).sort_values()
coef2.plot(kind='bar', title='Model Coefficients')
```


Model Report

RMSE : 1129

CV Score : Mean - 1130 | Std - 44.6 | Min - 1076 | Max - 1217

r2 score : 0.5625

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1c15913f10>
```



## Decision Tree Model

```
from sklearn.tree import DecisionTreeRegressor
DT = DecisionTreeRegressor(max_depth=15, min_samples_leaf=100)
modelfit(DT, train_df, test_df, predictors, target, IDcol, 'DT.csv')
```

```
coef3 = pd.Series(DT.feature_importances_, predictors).sort_values(ascending=False)
coef3.plot(kind='bar', title='Feature Importances')
```

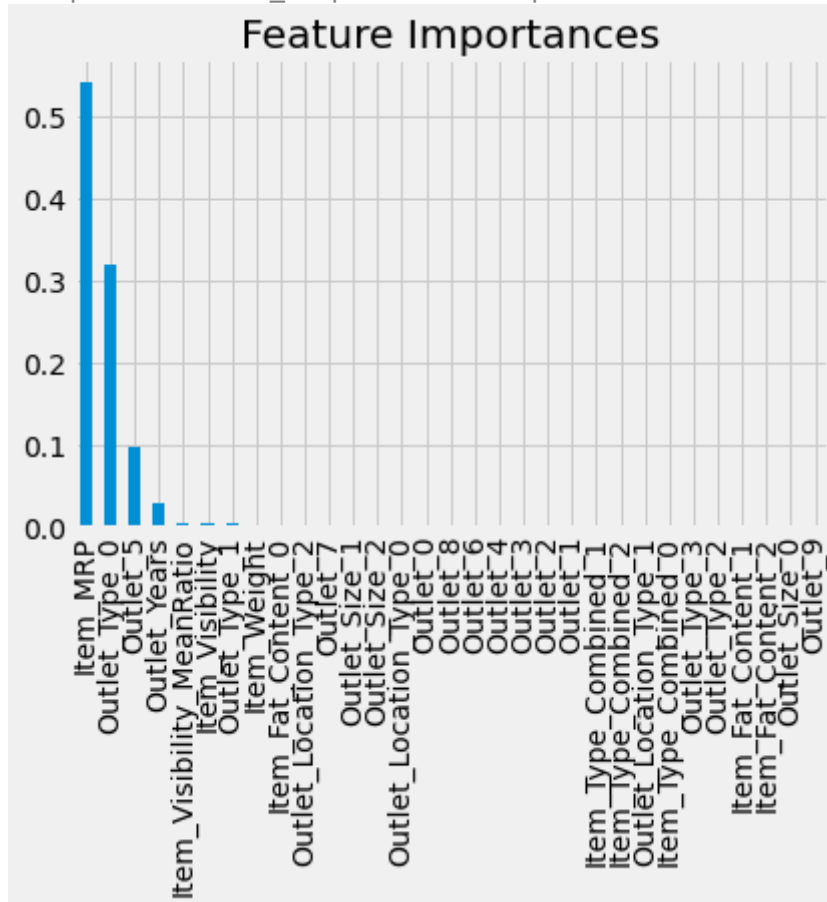
Model Report

RMSE : 1058

CV Score : Mean - 1091 | Std - 45.42 | Min - 1003 | Max - 1186

r2 score : 0.6158

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f1c15cb4750>



## Random Forest Model

```
RF = DecisionTreeRegressor(max_depth=8, min_samples_leaf=150)
```

```
modelfit(RF, train_df, test_df, predictors, target, IDcol, 'RF.csv')
```

```
coef4 = pd.Series(RF.feature_importances_, predictors).sort_values(ascending=False)
```

```
coef4.plot(kind='bar', title='Feature Importances')
```

Model Report

RMSE : 1069

CV Score : Mean - 1097 | Std - 43.41 | Min - 1028 | Max - 1180

r2 score : 0.6077

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f1c15b15150>



xgboost



```
from xgboost import XGBRegressor
```

```
my_model = XGBRegressor(n_estimators=1000, learning_rate=0.05)
my_model.fit(train_df[predictors], train_df[target], early_stopping_rounds=5,
              eval_set=[(test_df[predictors], test_df[target])], verbose=False)
```

```
[08:42:53] WARNING: /workspace/src/objective/regression_obj.cu:152: reg:linear is now
XGBRegressor(learning_rate=0.05, n_estimators=1000)
```



```
#Predict training set:
train_df_predictions = my_model.predict(train_df[predictors])
```

```
# make predictions
predictions = my_model.predict(test_df[predictors])
```

```
from sklearn.metrics import mean_absolute_error
print("Mean Absolute Error : " + str(mean_absolute_error(predictions, test_df[target])))
print("RMSE : %.4g" % np.sqrt(metrics.mean_squared_error((train_df[target]).values, train_
print("r2 score : %.4g" % metrics.r2_score((train_df[target]).values, train_df_predictions
```

Mean Absolute Error : 220.26657130725278

RMSE : 1052

r2 score : 0.6197

```
IDcol.append(target)
submission = pd.DataFrame({ x: test_df[x] for x in IDcol})
submission.to_csv("XGboost.csv", index=False)
```

